







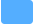
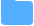
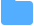









NLP-Project Public

  12 Branches  0 Tags  t [Go to file](#) [Add file](#) [About](#) [Code](#) ... 

 FaithWambugu [Delete sentiment_analysis2.ipynb](#) 89d884f · now 

	Data	Added folders	12 hours ago
	Images	Added Presentation and i...	7 hours ago
	Notebook PDF	Final Touches	9 minutes ago
	Presentation	Added folders	12 hours ago
	.gitignore	Configured gitignore to e...	last week
	README.md	Revise README for clarit...	11 hours ago
	environment.yml	Update environment.yml ...	2 days ago
	requirements.txt	Merge pull request #14 fr...	10 hours ago
	requirements_upd...	Add updated requiremen...	11 hours ago
	sentiment_analysis...	Add files via upload	3 minutes ago
	sentiment_app.py	Fix: update Streamlit app ...	14 hours ago

No description, website, or topics provided.

 [Readme](#)

 [Activity](#)

 0 stars

 0 watching

 1 fork

Releases

No releases published

[Create a new release](#)

Packages

No packages published

[Publish your first package](#)

Contributors 6



Languages

- Jupyter Notebook 99.8%
- Python 0.2%

Suggested workflows

Based on your tech stack



Python Package using Anaconda

[Configure](#)

Create and test a Python package on multiple Python versions using Anaconda for package management.



Publish Python Package

[Configure](#)

Publish a Python Package to PyPI on release.



SLSA Generic generator

[Configure](#)

Generate SLSA3 provenance for your existing release workflows

[More workflows](#)[Dismiss suggestions](#)[README](#)

Apple & Google Tweet Sentiment Analyzer

Real-Time Sentiment Monitoring for Tech Distributors

This project leverages **Transfer Learning** with the **DistilBERT** model to provide third-party technology distributors with critical, real-time insights into public sentiment towards Apple and Google products. The primary goal is to establish a robust alerting system by maximizing the detection of negative customer feedback.

Our final model achieved an exceptional 90 **Recall** on the Negative sentiment class, ensuring that the highest-risk tweets are flagged immediately for business action.



Business Context & Motivation

Third-party distributors operate in a high-stakes, competitive environment where product perception dictates sales and inventory strategy. They currently lack accessible tools for automated, large-scale sentiment analysis, leaving them vulnerable to market shifts and slow to react to customer dissatisfaction.

This solution bridges that gap by transforming unstructured social media data into **actionable intelligence**, enabling distributors to:

- **Prevent Missed Warnings:** Proactively address potential product issues or widespread customer frustration.
- **Optimize Inventory:** Align stocking decisions with real-time brand perception.
- **Inform Marketing:** Adjust campaigns and messaging based on evolving public opinion.

Key Technical Features

The solution employs a hybrid modeling approach to serve both alerting and strategic analysis needs:

1. Binary Classification (Alerting Core):

- **Goal:** High Recall for the Negative class (0.90 achieved).
- **Model:** DistilBERT fine-tuned on the Negative vs. Positive classes.

2. Multiclass Classification (Strategic Trends):

- **Goal:** Contextual insight into all sentiment types (Negative, Neutral, Positive).
- **Model:** DistilBERT fine-tuned on the three sentiment classes.

3. Advanced Handling of Imbalance:

- **Technique:** Combination of **Data Augmentation** (Synonym Replacement) and a custom **Focal Loss** function were implemented to effectively address the severe class imbalance, surpassing the performance of traditional SMOTE/Class Weights.



Performance Summary

The Deep Learning approach decisively outperformed the classical Machine Learning models (Logistic Regression, Random Forest, XGBoost) across all key metrics.

Model	Classification Task	Negative Recall	Macro F1-Score
DistilBERT	Binary (Alerting)	0.90	0.89
DistilBERT	Multiclass (Trends)	0.82	0.69
XGBoost	Binary (ML Champion)	0.57	0.73

Technologies & Libraries

- **Primary NLP Framework:** Hugging Face transformers (DistilBERT)
- **Deep Learning:** PyTorch (torch)
- **Machine Learning:** Scikit-learn, XGBoost, imblearn (SMOTE)
- **Feature Engineering:** TF-IDF, GloVe Word Embeddings
- **Visualization & Deployment:** Streamlit, Matplotlib, Seaborn
- **Experiment Tracking:** Weights & Biases (wandb)

How to Run the Project

The core project output is the interactive Streamlit dashboard.

1. Prerequisites

Ensure you have **Python 3.8+** and a package manager (pip or conda).

2. Clone the Repository

```
git clone git@github.com:FaithWambugu/NLP-Project.git
cd <repository_name>
```



3. Set Up the Environment

Install all required dependencies, including PyTorch and the Hugging Face libraries:

```
pip install -r requirements.txt
```



4. Launch the Streamlit Application

The application provides the final deployed solution, visualizing binary alerts and multiclass trends.

```
streamlit run sentiment_app.py
```



5. View the Full Analysis

All data cleaning, EDA, ML, and DL modeling steps are fully documented in the main Jupyter Notebook:

```
jupyter notebook
```



(Open the main notebook file and run the cells in sequence.)

Data Source & Structure

- **Source:** CrowdFlower Brands and Product Emotions dataset, available via `data.world`.
- **Total Records:** Over 9,000 tweets.
- **Key Columns:**
 - `tweet_text` : The raw text for analysis.

- `emotion_in_tweet_is_directed_at` : The specific brand or product.
- `is_there_an_emotion_directed_at_a_brand_or_product` : The original sentiment label (Positive, Negative, No emotion).

(Note: To complete the deployment, ensure your `requirements.txt` includes `torch` , `transformers` , `streamlit` , and `wandb` , and verify the paths to your saved DistilBERT models are correct in `sentiment_app.py`)